

Shot aggregating strategy for near-duplicate video retrieval

Vignesh Srinivasan, Frederic Lefebvre, Alexey Ozerov

► To cite this version:

Vignesh Srinivasan, Frederic Lefebvre, Alexey Ozerov. Shot aggregating strategy for near-duplicate video retrieval. 23rd European Signal Processing Conference (EUSIPCO 2015), Aug 2015, Nice, France. hal-01164350

HAL Id: hal-01164350

<https://hal.inria.fr/hal-01164350>

Submitted on 16 Jun 2015

HAL is a multi-disciplinary open access archive for the deposit and dissemination of scientific research documents, whether they are published or not. The documents may come from teaching and research institutions in France or abroad, or from public or private research centers.

L'archive ouverte pluridisciplinaire **HAL**, est destinée au dépôt et à la diffusion de documents scientifiques de niveau recherche, publiés ou non, émanant des établissements d'enseignement et de recherche français ou étrangers, des laboratoires publics ou privés.

SHOT AGGREGATING STRATEGY FOR NEAR-DUPLICATE VIDEO RETRIEVAL

Vignesh Srinivasan, Frederic Lefebvre and Alexey Ozerov

Technicolor

975 avenue des Champs Blancs,

CS 17616, 35576 Cesson Sévigné, France

{vignesh.srinivasan, frederic.lefebvre, alexey.ozerov}@technicolor.com

ABSTRACT

In this paper, we propose a new strategy for near-duplicate video retrieval that is based on shot aggregation. We investigate different methods for shot aggregation with the main objective to solve the difficult trade-off between *performance*, *scalability* and *speed*. The proposed shot aggregation is based on two steps. The first step consists of keyframes selection. And the second one is the aggregation of the keyframes per shot. The aggregation is performed by applying Fisher vector on the descriptors computed on the selected keyframes. We demonstrate that the *scalability* and the *speed* are tackled by a sparse video analysis approach (i.e. extracting only few keyframes) combined with shot aggregation, while the *performance* is discussed around the choice of the aggregation strategy. The *performance* is evaluated on the CC_WEB_VIDEO dataset that is designed for the near-duplicate video retrieval assessment and for which some experiments have been conducted by different authors.

1. INTRODUCTION

Over the last decade, several web services offer online-storage for personal video backup, such as Dropbox, OneDrive, or for video sharing such as Youtube, Dailymotion. The problem of efficient storage, retrieval, and copyright infringement has motivated the online storage providers to develop the concept of near-duplicates in their platforms. Near-duplicate video content (NDVC) is often defined as identical or approximately identical videos but different in file formats, encoding parameters, photometric variations, editing operations, lengths, and certain modifications [1] such as variations in camera viewpoint, setting [2] and camera motion [3]. Cherubini *et al.* [4] in agreement with [3] articulates the definition of near-duplicate video clip from a user's perspective by taking into account the participants activity on video sharing websites. Jiang *et al.* [5] defines NDVC as the same scenes originally captured from two different cameras. Since there is no unique consensus on the NDVC definition in the literature, particular approaches rely on some datasets, where it is annotated whether two videos are near-duplicate or not. In this work we will follow this data-driven definition.

There are a number of datasets that challenge video fingerprint and give data-driven definitions of near-duplicate videos. The Muscle-VCD dataset [6] and the TRECVID dataset provided by the National Institute of Science and Technology [7] are essentially for video copy detection. The CC_WEB_VIDEO dataset [1] does not consider the large scale but proposes NDVC content from real world rather than artificially created content. UQ_VIDEO [2] is

a combined dataset created by injecting more videos to the existing CC_WEB_VIDEO. VCDB is a new dataset [5] for partial copy detection in videos containing about 100,000 videos downloaded from the internet. The evaluation procedure needs to compare the scores obtained by our algorithm with state of the art in the same test collection. CC_WEB_VIDEO dataset benefits to some relevant comparison studies [8], [9] in near duplicate use case. For this reason, we have selected CC_WEB_VIDEO as the reference dataset for our experiments.

To achieve a good trade-off between *performance*, *scalability* and *speed*, most of the video fingerprint algorithms manage the video as a succession of frames. For each selected frame, some attributes or descriptors are extracted in a sparse or dense manner. The quality of attributes affects the accuracy of a near-duplicate video retrieval (NDVR) system while their size and number impact the scalability and the speed of retrieval. Color histograms [1] used as a global video signature are computationally expensive and also fail in case of photometric variations in the videos. An improvement over this issue was proposed [1, 2], which consists in using a hierarchical approach employing local signatures. Shang *et al.* [9] introduced a video signature method based on a binary spatio-temporal feature. A temporal approach is also elaborated in [10] by finding temporal relations among patterns. To improve scalability and the speed of retrieval, a temporal sparse approach consisting of detecting keyframes is employed [11]. Ordinal relations are extracted only on these keyframes using conditional entropy and local binary pattern methods. However, this method is not robust to spatial editing and hence brings the performance down. Visual descriptors aggregation methods [12–20] became more and more popular because of their robustness to different transformations or editing in the videos. Aggregated descriptors tackle the sparsity of the image descriptor but do not handle the temporal aspect. Note also that selecting “relevant” keyframes is a way to improve the scalability and speed in video description. Uniformly sampled keyframes are used in [9, 16, 18, 20, 21] which leads to an extra storage cost because of redundant information. One keyframe per shot (the center of the shot) is extracted in [1, 13, 15]. Such a method of keyframe extraction leads to a loss of information over the entire shot and hence to a performance drop, especially in the case of long shots. To improve the accuracy, a method consists in increasing the number of features [22] but this leads also to an extra cost storage.

The aim of our proposal is to maintain a high performance while increasing the scalability and speed of retrieval. First, the keyframes are selected using the method described in [23, 24]. Shot boundaries are found in a video and stable keyframes are extracted from inside these shots in a non-uniform manner. The scale-invariant feature transform (SIFT) [25] descriptors are extracted from each of the keyframes. A feature vector is calculated per shot by aggregating

This work was partially supported by the FP7 European integrated project AXES. <http://www.axes-project.eu/>

all extracted descriptors in a shot in a single Fisher vector, which is a simplification of the Fisher kernel [26]. Finally, the retrieved videos are ranked using two different strategies: a voting strategy or a hidden Markov model (HMM)-based strategy that allows exploiting temporal coherence between sequences of shots. The main contribution of this paper is to combine a sparse video analysis approach, i.e., selecting just few keyframes, with different aggregation methods that should lead a good trade off between *performance*, *scalability* and *speed*.

The rest of the paper is organized as follows. Section 2 describes our proposed shot aggregation method in detail. Section 3 is devoted to experiments. Conclusions are drawn in section 4.

2. PROPOSED METHOD

A general scheme of the proposed method is represented in Fig. 2. In the following subsections the method is described step by step.




Fig. 1. General scheme of the proposed method.

2.1. Keyframe Selection

The goal of this step is to find the best frames in a video to fingerprint with a reasonable density rate. A shot being a temporal section where the video activity shall be constant, the method we use [23,24] consists in finding shot boundaries and the best stable frames in each shot. Shot boundaries are the two frames that surround the shot while the best stable frames are the frames with the smallest content variation along the shot. The activity is captured by analysing a perceptual distance between successive frames [24]. This perceptual distance is defined as the Euclidean distance between perceptual hash (a global descriptor) of neighbour frames. Perceptual hash we used is a so-called Radon soft hash algorithm (RASH) [27] that is computed on radial strip (set of points on a line passing through the image center), with orientation $\theta \in [1^\circ, 180^\circ]$. More details on shot detection can be found in [24].

We have observed that selecting only the best stable frame is relevant in case of short temporal section but it is not enough if the shot is longer than few seconds. For this reason we have decided to extend the number of stable frames to a maximum of A stable frames per shot under the constraint that the temporal distance between two stable frames is at least B frames. Values of A and B drive the density of stable frames in a temporal section. We have chosen $A = 10$ and $B = 20$. We have also observed that the stable frames are usually the best in the shot in terms of blurriness (due to motion), darkness or brightness but some artefacts due to image macro block compression persist. For each stable frame detected, we detect if it is a I, P, or B-type frame. If the stable frame is not an I frame, we search in a neighbour window (of up to 5 frames) if an I frame is available. I frames are better encoded than P or B frames and generate less of compression artefacts, so it is better to apply image fingerprint on I frames in the latter stage.

Once the keyframes are selected, they have to be described. SIFT local descriptors aggregated by Fisher vector techniques are dedicated to this task.

2.2. Fisher Vector Encoding

The Fisher kernel [19] combines the advantages of methods based on generative statistical models and those of discriminative methods. The Fisher vector is the normalized gradient of the log-likelihood of the data sample with respect to the model parameters, the model being pre-trained from some training data.¹ The Fisher vector can be derived as a special case of the Fisher kernel [26]. The Fisher vector is an image representation obtained by pooling local image features, and it results in a dense and compact image representation that was shown more efficient than the standard bag of visual words (BOVW) for image classification and categorization [26,28].

Consider a D -dimensional feature vector $X = [x_1, x_2, \dots, x_D]^T$ (here we use dense SIFT) extracted from an image and the parameters of the Gaussian mixture model (GMM) to be $\theta = \{\mu_k, \Sigma_k, \pi_k\}_{k=1}^K$, where μ_k , Σ_k and π_k denote respectively the mean vector, the covariance matrix and the weight of k -th component. The GMM associates each feature vector X to a component k in the mixture with a strength given by the posterior probability. Thus, the assignment to a given GMM component is done in a soft manner, which is a fundamental difference between the Fisher vector and the BOVW model that is based on a hard assignment. We here compute the Fisher vector with respect to GMM means μ_k only [29], and thus the Fisher vector length is

$$L = D \times K. \quad (1)$$

For example for a SIFT descriptor of size 128 and a GMM with 64 components, we get a Fisher vector of length - $128 \times 64 = 8192$.

Unlike the previous works [9, 16, 18, 20, 21], where either only one frame per shot or uniformly sampled frames were considered, our method considers several non-uniformly selected stable frames, aka keyframes, per shot. By not considering uniformly sampled frames, our method eliminates redundant information and this should help in its scalability. Once dense SIFT and Fisher vector are applied on keyframes, the shot aggregation of the local features is performed. We consider two alternative ways for that:

1. Shot level SIFT aggregation (**S-AGG**) : Dense SIFT features computed for each of the keyframes in a shot are aggregated into a single Fisher vector per shot:

$$FV_{shot} = \mathcal{G} \left(\left\{ \{X_{m,n}\}_{n=1}^{N_k(m)} \right\}_{m=1}^{M_s} \right), \quad (2)$$

where $X_{m,n}$ is the n -th SIFT vector in the m -th keyframe, $N_k(m)$ is the total number of SIFT vectors in the m -th keyframe, M_s is the total number of keyframes in the shot, and $\mathcal{G}(\cdot)$ is the Fisher vector aggregation function.

2. Shot level Fisher averaging (**F-AV**): For each keyframe one Fisher vector is computed from the corresponding dense SIFT features. Then, Fisher vectors belonging to a single shot are averaged to obtain a single Fisher vector for this shot:

$$FV_{shot} = \frac{1}{M_s} \sum_{m=1}^{M_s} \mathcal{G} \left(\{X_{m,n}\}_{n=1}^{N_k(m)} \right). \quad (3)$$

¹In case of image representation this training data may be a set of descriptors extracted from a big set of images.

2.3. Ranking

To rank the retrieved videos, we propose two strategies: the voting strategy and the HMM-based strategy. Both strategies are using a similarity measure that is the Euclidean distance between the descriptors, i.e., Fisher vectors, of the query video and the descriptors of all the videos in the dataset. Let $\mathcal{D} = [d_{ij}]_{i,j=1}^{I,J}$ matrix of similarity measures between the query video and a video from the reference dataset, where d_{ij} is the Euclidean distance between the i^{th} descriptor of the query and the j^{th} descriptor of the reference dataset, indices i and j enumerate to either the shots or the keyframes depending on the method, and I and J are the corresponding total numbers of shots or keyframes.

2.3.1. Majority voting

We rely on the voting strategy introduced in [30]. Given \mathcal{D} the voting similarity between the query video and a video from the reference dataset is computed as

$$V = \sum_{i=1}^I \left(\max_{k=1,\dots,K} [\tilde{d}_{ik}] - \min_{j=1,\dots,J} [d_{ij}] \right), \quad (4)$$

where $\tilde{\mathcal{D}} = [\tilde{d}_{ik}]_{i,k=1}^{I,K}$ is the matrix of similarity measures between the query video and all videos from the reference dataset, with index $k = 1, \dots, K$ enumerating over all videos the shots or the keyframes depending on the method. The videos can now be ranked according to the votes received.

2.3.2. HMM-based strategy

Voting similarity described in the previous section does not take into account temporal consistency between different keyframes or shots. However, in the near-duplicate content the keyframes or shots keep usually the same order in the query and the reference videos, and at the same time this does not happen in 100 %, since some shot swaps, insertions or deletions are possible as well. Here we would like exploiting such a temporal consistency, while having an approach that tolerates temporally non-consistent exceptions. For that we are using an HMM-based strategy that allows taking into account temporal consistency via a most likely path decoded within similarity matrix \mathcal{D} that usually relies distances d_{ij} in a monotone manner (see red path on Fig. 2), while non-monotone behaviour is possible as well. HMM-based strategy was also used to detect copied segments in [31].

This is achieved as follows. Given a similarity matrix \mathcal{D} between a query video and one reference video, we consider an I -length observation sequence decoded within a J -state HMM. The similarity measure between these two videos is now computed as the log-likelihood of the most likely sequence of states $q = \{q_i\}_{i=1}^I$ ($q_i \in \{1, \dots, J\}$). In other words

$$V = \max_q \log p(q), \quad (5)$$

where

$$\log p(q) = -c \times \sum_{i=1}^I d_{iq_i} + \sum_{i=2}^I \log p(q_i | q_{i-1}) + \log p(q_1), \quad (6)$$

is a log-likelihood of a sequence q that is defined by replacing the observation log-likelihood at time i and state j by $c d_{ij}$ with c being a positive constant (here we use $c = 0.001$), initial probabilities $p(j)$

are fixed to equal values and transition probabilities $p(j|j')$ are defined so as to favour short state transitions forward, while not forbidding any other state transitions, e.g., backward or far forward (and example of transition probabilities values can be found on Fig. 2). This allows temporal consistency check, while being tolerant to exceptional cases where this consistency does not hold (e.g., due to video editing).

This maximum log-likelihood (5) can be efficiently computed by dynamic programming or Viterbi algorithm [32] relying on forward and backward propagation. Since we only need the maximum log-likelihood value and do not really need the most likely state sequence q , we are using forward propagation only.




Fig. 2. HMM-based strategy

3. EXPERIMENTS

Here we evaluate our methods on the CC.WEB.VIDEO dataset, since it is a popular benchmark for the NDVR task and results of state-of-the-art methods for this benchmark are available [8, 9]. In this dataset 24 different queries are issued from YouTube, Google Video, and Yahoo! and the corresponding search results are collected to form a dataset consisting of 12790 videos which is split into 24 subsets based on the queries. We pre-trained a GMM for the Fisher vector computation using a completely different dataset, built with stable frames collected from thirty hours of videos, and for varying number of components: 64, 128, 256. The keyframes are then densely sampled and SIFT features are computed on each patch. Fisher vectors are then computed based on the two shot aggregation strategies, **S_AGG** and **F_AV**, described in section 2.2. To evaluate the performance of our methods, we have two additional experimental set-ups, **A_KF** and **C_SH**:

- **A_KF**: Fisher vectors for all the M_s keyframes from a shot are computed and no shot aggregation is performed.
- **C_SH**: We extract one middle keyframe per shot, similar to [1], but using our shot detection method described in section 2.1. Fisher vectors are computed for all the keyframes.

We also consider two baseline state-of-the-art methods [8, 9]:

- **B_CE**: In [9], keyframes are sampled uniformly every second and a conditional entropy method is applied to all the keyframes.
- **B_CCA**: In [8], a canonical correlation analysis (CCA) is performed between two videos to check for similarity. The CCA is applied on the features extracted from the videos.

For all the methods based on the Fisher vectors, we perform tests with 3 different GMM components: 64, 128, 256. The lower the number of components, the smaller is the Fisher vector size and hence the storage requirements, which leads in turn to a better scalability.




Fig. 3. Averaged precision-recall curves across 24 queries for our methods.

3.1. Results

We evaluated the performances for each of the set-up, and the corresponding precision-recall curves [8, 9] are plotted on Figure 3. The best overall performance of **F_AV** method is obtained with 128 GMM components. **S_AGG-128** performs similar to **F_AV-128**, while **A_KF-128** has a slightly higher precision for lower recall values. However, as recall is increased, the precision drops heavily. Since for **C_SH-128** method only the middle keyframes from the shot are considered, the area under the precision-recall curve is relatively smaller than for other methods considered here.

Figure 4 plots the comparison of our methods **F_AV-128** and **F_AV-128-HMM** with the baseline state-of-the-art methods **B_CE** and **B_CCA**.² We can see that our method **F_AV-128** has a performance similar to **B_CCA**, while **F_AV-128-HMM** has a higher precision than **B_CCA**. This is due to the fact that the HMM-based strategy checks the temporal coherence between matched shots or keyframes sequences using a probabilistic model. HMM-based strategy does not improve the recall but reduces the false positive rate. Both **F_AV-128** and **F_AV-128-HMM** outperforms **B_CE** for lower recall values.

3.2. Storage and Complexity

Our shot aggregation methods require an average 0.210 MB for GMM with 64 components and the storage requirements increase linearly with respect to the GMM size. **F_AV-128** requires a higher storage than **F_AV-64** but performs better compared to it. Since descriptors are computed and stored for many keyframes in each shot, **A_KF** requires the maximum storage space among our methods, thus affecting its scalability. The computational expense is also increased compared to **F_AV-128**. Keyframes are sampled uniformly in **B_CE** and hence it requires more storage compared to **F_AV-128**. With relatively smaller storage requirements and a less expensive computation, **F_AV-128** provides a similar performance to **B_CCA**. **F_AV-128-HMM** is computationally more expensive than **F_AV-128** but it improves the performance.

²Precision-recall curves of the baseline methods are taken from the corresponding papers [8, 9].




Fig. 4. Comparison of our best methods with two baselines.

4. CONCLUSION

This paper introduced several shot aggregation strategies for NDVR. The keyframes from the videos are extracted using a shot boundary detector followed by non-uniform stable frame selection from the shots. Shot aggregation methods are applied to each shot in a video from which one single Fisher vector per shot is computed.

This strategy provides similar or higher performance than the baseline methods that are based on single keyframe extraction from the center of each shot. It also provided better scalability compared to baseline state-of-the-art methods that are based on uniform keyframes sampling. The combination of shot aggregation strategy and an HMM-based strategy for ranking the near-duplicates improved the performance compared to the voting strategy at the expense of a slightly higher computational load. Our best method is robust to temporal and spatial editing, photometric and geometric variations and gives a similar performance compared to the baseline state-of-the-art methods [8, 9] with smaller storage requirements. It also provides a good trade-off between *performance*, *scalability* and *speed*.

REFERENCES

- [1] Xiao Wu, Alexander G Hauptmann, and Chong-Wah Ngo, "Practical elimination of near-duplicates from web video search," in *Proceedings of the 15th international conference on Multimedia*. ACM, 2007, pp. 218–227.
- [2] Heng Tao Shen, Xiaofang Zhou, Zi Huang, Jie Shao, and Xiangmin Zhou, "Uqlips: a real-time near-duplicate video clip detection system," in *Proceedings of the 33rd international conference on Very large data bases*. VLDB Endowment, 2007, pp. 1374–1377.
- [3] Arslan Basharat, Yun Zhai, and Mubarak Shah, "Content based video matching using spatiotemporal volumes," *Computer Vision and Image Understanding*, vol. 110, no. 3, pp. 360–377, 2008.
- [4] Mauro Cherubini, Rodrigo De Oliveira, and Nuria Oliver, "Understanding near-duplicate videos: a user-centric approach," in *Proceedings of the 17th ACM international conference on Multimedia*. ACM, 2009, pp. 35–44.

- [5] Yu-Gang Jiang, Yudong Jiang, and Jiajun Wang, "VCDB: A Large-Scale Database For Partial Copy Detection In Videos," in *Computer Vision-ECCV 2014*, pp. 357–371. Springer, 2014.
- [6] J Law-To, A Joly, and N Boujemaa, "Muscle-VCD-2007: a live benchmark for video copy detection," 2007.
- [7] "U.S National Institute Of Science And Technology: TRECVID video retrieval evaluation. <http://trecvid.nist.gov/>," .
- [8] Jiajun Liu, Zi Huang, Heng Tao Shen, and Bin Cui, "Correlation-based retrieval for heavily changed near-duplicate videos," *ACM Transactions on Information Systems (TOIS)*, vol. 29, no. 4, pp. 21, 2011.
- [9] Lifeng Shang, Linjun Yang, Fei Wang, Kwok-Ping Chan, and Xian-Sheng Hua, "Real-time large scale near-duplicate web video retrieval," in *Proceedings of the international conference on Multimedia*. ACM, 2010, pp. 531–540.
- [10] Chun-Chieh Hsu Chien-Peng Ho Chien-Li Chou, Hua-Tsung Chen and Suh-Yin Lee, "Near-duplicate video retrieval by using pattern-based prefix tree and temporal relation forest," in *International Conference on Multimedia and Expo (ICME)*. IEEE, 2014, vol. 1,6, pp. 14–18.
- [11] A.A. Tonge S.D Thepade, "An optimized key frame extraction for detection of near duplicates in content based video retrieval," in *International Conference on Communications and Signal Processing (ICCS)*. IEEE, 2014, pp. 1087–1091.
- [12] Yu-Gang Jiang, Chong-Wah Ngo, and Jun Yang, "Towards optimal bag-of-features for object categorization and semantic video retrieval," in *Proceedings of the 6th ACM international conference on Image and video retrieval*. ACM, 2007, pp. 494–501.
- [13] Yu-Gang Jiang and Chong-Wah Ngo, "Visual word proximity and linguistics for semantic video indexing and near-duplicate retrieval," *Computer Vision and Image Understanding*, vol. 113, no. 3, pp. 405–414, 2009.
- [14] Wan-Lei Zhao, Song Tan, and Chong-Wah Ngo, "Large-scale near-duplicate web video search: challenge and opportunity," in *Multimedia and Expo, 2009. ICME 2009. IEEE International Conference on*. IEEE, 2009, pp. 1624–1627.
- [15] Wan-Lei Zhao, Xiao Wu, and Chong-Wah Ngo, "On the annotation of web videos by efficient near-duplicate search," *Multimedia, IEEE Transactions on*, vol. 12, no. 5, pp. 448–461, 2010.
- [16] Matthijs Douze, Hervé Jégou, and Cordelia Schmid, "An image-based approach to video copy detection with spatio-temporal post-filtering," *Multimedia, IEEE Transactions on*, vol. 12, no. 4, pp. 257–266, 2010.
- [17] Matthijs Douze, Hervé Jégou, Cordelia Schmid, and Patrick Pérez, "Compact video description for copy detection with precise temporal alignment," in *Computer Vision-ECCV 2010*, pp. 522–535. Springer, 2010.
- [18] Mohamed Ayari, Jonathan Delhumeau, Matthijs Douze, Hervé Jégou, Danila Potapov, Jérôme Revaud, Cordelia Schmid, Jiangbo Yuan, et al., "Inria@ trecvid'2011: Copy detection & multimedia event detection," in *TRECVID*, 2011.
- [19] Florent Perronnin and Christopher Dance, "Fisher kernels on visual vocabularies for image categorization," in *Computer Vision and Pattern Recognition, 2007. CVPR'07. IEEE Conference on*. IEEE, 2007, pp. 1–8.
- [20] Hervé Jégou, Matthijs Douze, Cordelia Schmid, and Patrick Pérez, "Aggregating local descriptors into a compact image representation," in *Computer Vision and Pattern Recognition (CVPR), 2010 IEEE Conference on*. IEEE, 2010, pp. 3304–3311.
- [21] Apostol Natsev, John R Smith, Matthew Hill, Gang Hua, Bert Huang, Michele Merler, Lexing Xie, Hua Ouyang, and Mingyuan Zhou, "Ibm research trecvid-2010 video copy detection and multimedia event detection system," in *Proceedings of NIST TRECVID, Workshop*, 2010.
- [22] Zi Huang Heng Tao Shen Jingkuan Song, Yi Yang and Jiebo Luo, "Effective multiple feature hashing for large-scale near-duplicate video retrieval," *Multimedia, IEEE Transactions on*, vol. 15, no. 8, pp. 1997–2008, 2013.
- [23] Ayoub Massoudi, Frédéric Lefebvre, Claire-Hélène Demarty, Lionel Oisel, and Bertrand Chupeau, "A video fingerprint based on visual digest and local fingerprints," in *Image Processing, 2006 IEEE International Conference on*. IEEE, 2006, pp. 2297–2300.
- [24] Cedric De Roover, Christophe De Vleeschouwer, Frédéric Lefebvre, and Benoit Macq, "Robust video hashing based on radial projections of key frames," *Signal Processing, IEEE Transactions on*, vol. 53, no. 10, pp. 4020–4037, 2005.
- [25] David G. Lowe, "Distinctive image features from scale-invariant keypoints," *IJCV*, vol. 60(2), pp. 911–10, 2004.
- [26] Florent Perronnin, Jorge Sánchez, and Thomas Mensink, "Improving the fisher kernel for large-scale image classification," in *Computer Vision-ECCV 2010*, pp. 143–156. Springer, 2010.
- [27] Frédéric Lefebvre, Benoit Macq, Jean-Didier Legat, et al., "Rash: Radon soft hash algorithm," in *Proceedings of European Signal Processing Conference*. Toulouse, France, 2002, pp. 299–302.
- [28] Jorge Sánchez, Florent Perronnin, Thomas Mensink, and Jakob Verbeek, "Image classification with the fisher vector: Theory and practice," *International journal of computer vision*, vol. 105, no. 3, pp. 222–245, 2013.
- [29] Florent Perronnin, Yan Liu, Jorge Sánchez, and Hervé Poirier, "Large-scale image retrieval with compressed fisher vectors," in *Computer Vision and Pattern Recognition (CVPR), 2010 IEEE Conference on*. IEEE, 2010, pp. 3384–3391.
- [30] Hervé Jégou, Matthijs Douze, Cordelia Schmid, et al., "Exploiting descriptor distances for precise image search," *Research report, INRIA*, 2011.
- [31] Shikui Wei, Yao Zhao, Ce Zhu, Changsheng Xu, and Zhenfeng Zhu, "Frame fusion for video copy detection," *Circuits and Systems for Video Technology, IEEE Transactions on*, vol. 21, no. 1, pp. 15–28, 2011.
- [32] Lawrence Rabiner, "A tutorial on hidden markov models and selected applications in speech recognition," *Proceedings of the IEEE*, vol. 77, no. 2, pp. 257–286, 1989.